

Анализ эмоциональной окраски текста

на примере решения двух задач:

- 1) "Hate Speech and Offensive Language Dataset"
- 2) "Russian Troll Tweets"



Содержание

- 1 Постановка задачи
- 2 Ход решения
- 3 Особенности каждой задачи
- 4 Используемые библиотеки, алгоритмы, метрики
- 5 Итоги решений
- 6 Выводы
- 7 Список источников

Постановка задач

1

Постановка задач

Применение методов обработки текста для анализа тональности твитов в качестве задачи классификации, на примере двух задач:


1

"Hate Speech and Offensive Language Dataset - research hate-speech detection" (перевод с англ. яз. «Набор данных о высказываниях с ненавистью и оскорбительном языке — исследование обнаружения оскорбительных высказываний») - в дальнейшем сокращённо **HATE**

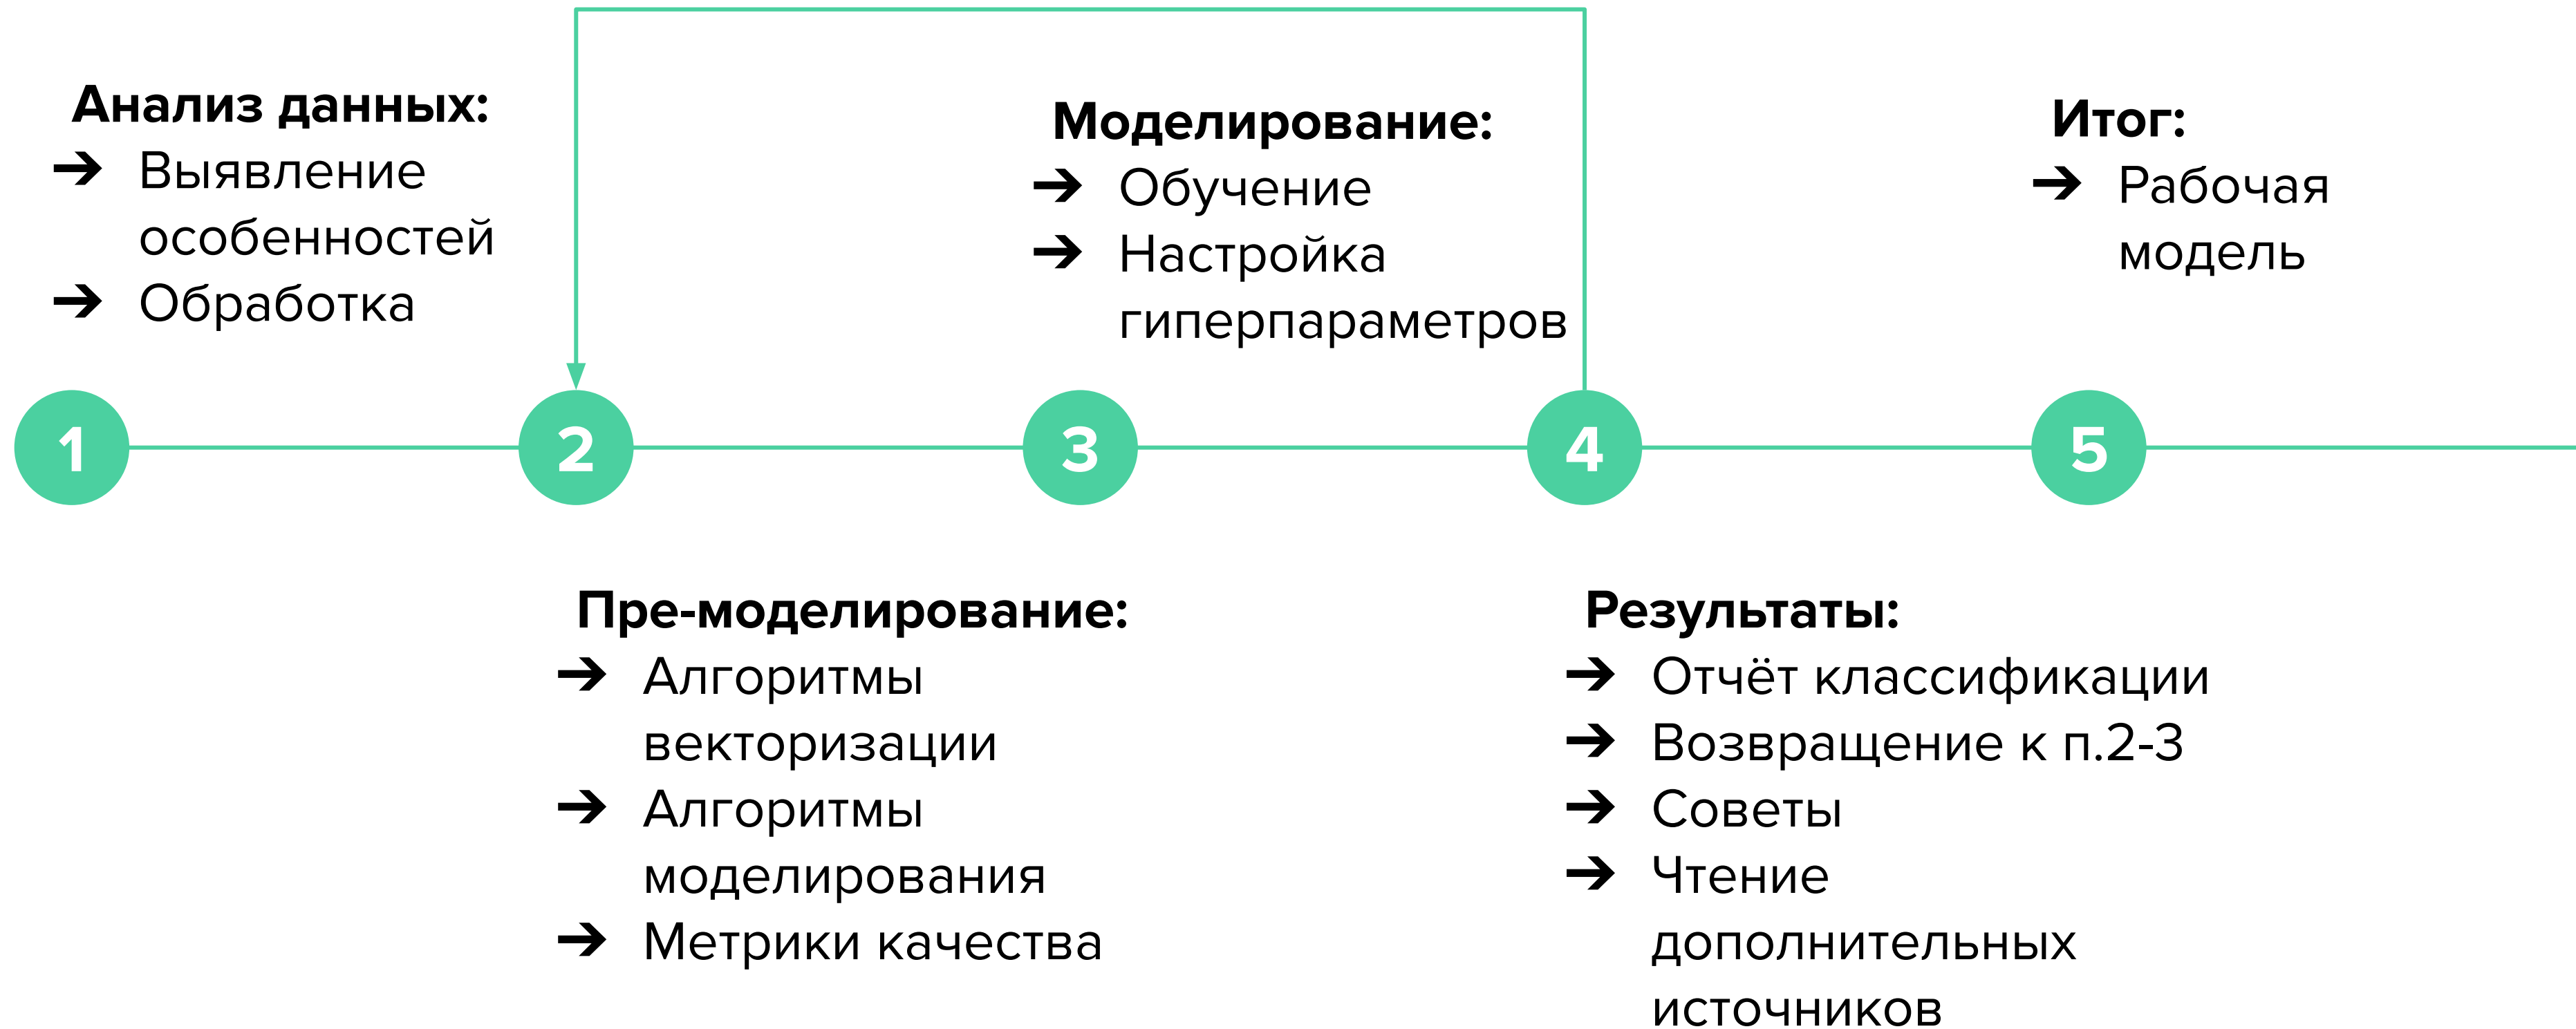
2

"Russian Troll Tweets - 3 million tweets from accounts associated with the 'Internet Research Agency'" («Твиты русских троллей — 3 миллиона твитов с аккаунтов, связанных с «Агентством интернет-исследований») - в дальнейшем сокращённо **TROLL**


Ход решения



Ход решения



Особенности каждой задачи



Особенности каждой задачи

Датасеты

Задача 1 - HATE

	count	hate_speech	offensive_language	neither	class	tweet
0	3	0	0	3	2	!!! RT @mayasolovely: As a woman you shouldn't...
1	3	0	3	0	1	!!!! RT @mleew17: boy dats cold...tyga dwn ba...
2	3	0	3	0	1	!!!!!! RT @UrKindOfBrand Dawg!!!! RT @80sbaby...
3	3	0	2	1	1	!!!!!! RT @C_G_Anderson: @viva_based she lo...
4	6	0	6	0	1	!!!!!! RT @ShenikaRoberts: The shit you...

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24783 entries, 0 to 25296
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0    count           24783 non-null  int64
1    hate_speech     24783 non-null  int64
2    offensive_language 24783 non-null  int64
3    neither         24783 non-null  int64
4    class           24783 non-null  int64
5    tweet           24783 non-null  object
dtypes: int64(5), object(1)
memory usage: 1.3+ MB
```

Задача 2 - TROLL

	external_author_id	author	content	region	language	publish_date	harvested_date	fol
0	2.385425e+09	MARRINABEREZKA	Обама принял решение по санкциям против Ирана ...	United States	Russian	11/11/2015 6:33	11/11/2015 6:34	26
1	2.534361e+09	ANETTANOVGOROD	Встреча Лаврова и Керри стартовала в Нью-Йорке...	Azerbaijan	Russian	9/27/2015 15:11	9/27/2015 15:11	16
2	1.612107e+09	LILJORDAMN	#IndieAdvancement Slim The Phenom @therealslim...	United States	English	12/3/2016 22:36	12/3/2016 22:36	60
3	3.254274e+09	FINDDIET	'@ozzyacaceres ozzy @laurengodfreyx1 Lauren @dj...	United States	English	8/5/2015 17:39	8/5/2015 17:39	3
4	1.647457e+09	COLINSNEVERLAND	This, BTW is why I don't instantly dismiss the...	United States	English	1/6/2016 18:02	1/6/2016 18:02	36

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2973371 entries, 0 to 2973370
Data columns (total 15 columns):
#   Column          Dtype
---  -
0    external_author_id float64
1    author           object
2    content          object
3    region           object
4    language         object
5    publish_date     object
6    harvested_date   object
7    following        int64
8    followers        int64
9    updates          int64
10   post_type        object
11   account_type     object
12   new_june_2018    int64
13   retweet          int64
14   account_category object
dtypes: float64(1), int64(5), object(9)
memory usage: 340.3+ MB
```

harvested_date	following	followers	updates	post_type	account_type	new_june_2018	retweet	account_category
11/11/2015 6:34	266	314	4160	RETWEET	Russian	1	1	NonEnglish
9/27/2015 15:11	166	153	1900	RETWEET	Russian	1	1	NonEnglish
12/3/2016 22:36	602	706	2531	RETWEET	left	0	1	LeftTroll
8/5/2015 17:39	3	200	21960	NaN	Commercial	1	0	Commercial
1/6/2016 18:02	364	202	127	RETWEET	Right	0	1	RightTroll

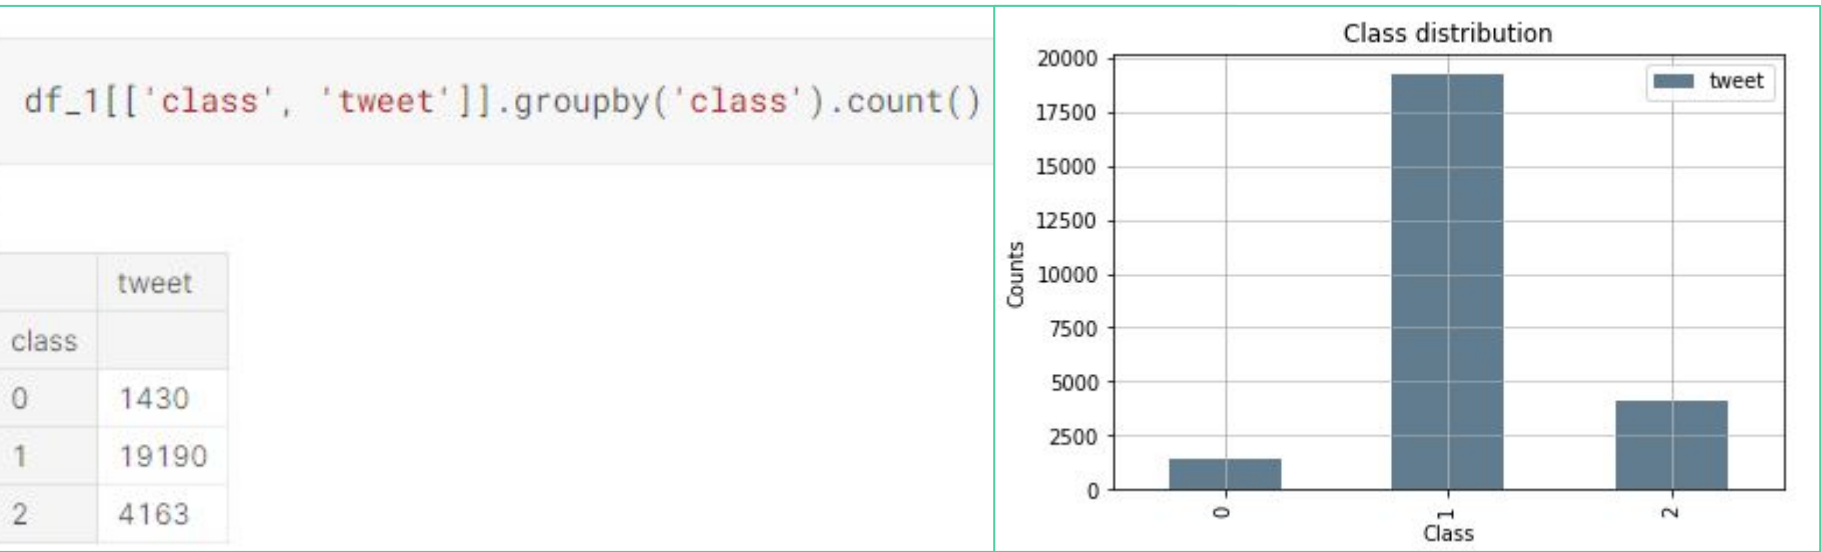
Особенности каждой задачи

Анализ

Задача 1 - HATE

```
df_1.tweet.iloc[0]
```

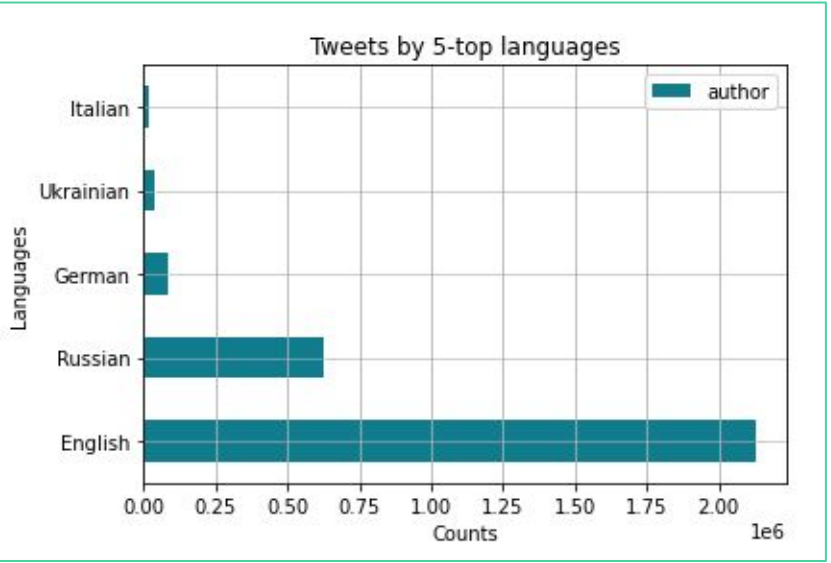
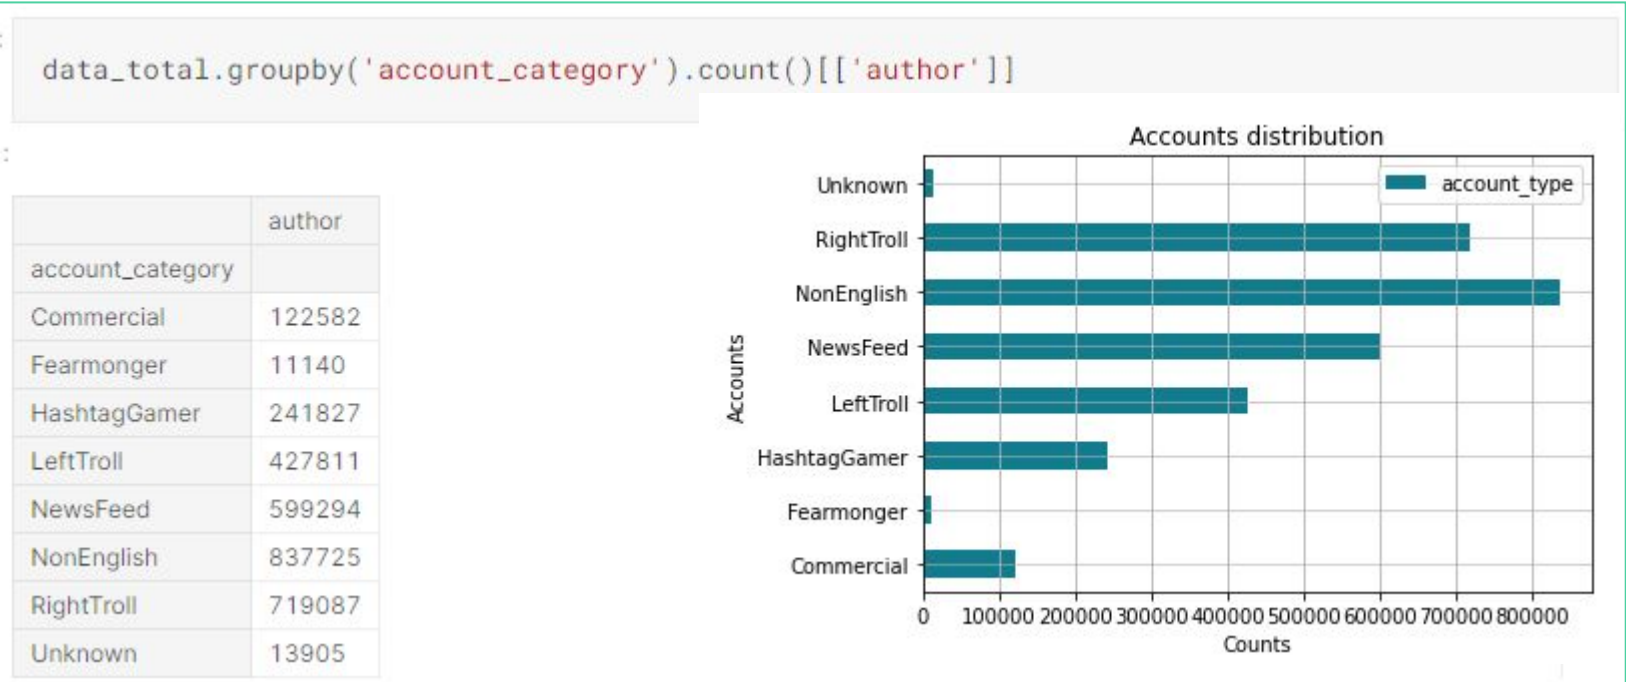
"!!! RT @mayasolovely: As a woman you shouldn't complain about cleaning up your house. & as a man you should always take the trash out..."



Задача 2 - TROLL


```
data_total[data_total['account_category'] == 'Commercial']['content'].iloc[0]
```

"@ozzyacaceras ozzy @laurengodfreyx1 Lauren @djhawes Danners @karleighwoelmer Karleigh h
ttp://t.co/QeKnVmkfxw https://t.co/ae3ItmkKin"



1. Векторизация
2. Регистр
3. Ники пользователей
4. Нехорошие ссылки
5. Слова с цифрами
6. Пунктуация
7. Стоп-слова
8. Лемматизация
9. Множественные пробелы

Используемые библиотеки, алгоритмы, метрики



Используемые библиотеки, алгоритмы, метрики

Библиотеки

Задача 1 - HATE

Загрузка данных

- os

Анализ

- numpy
- pandas
- matplotlib.pyplot
- tqdm

Предобработка

- re
- string.punctuation
- nltk.corpus.stopwords
- spacy
- sklearn.feature_extraction.text.CountVectorizer
- sklearn.feature_extraction.text.TfidfVectorizer
- nltk.word_tokenize
- imblearn.over_sampling.SMOTE
- sklearn.preprocessing.LabelEncoder

Метрики

- sklearn.metrics.classification_report

Задача 2 - TROLL

Моделирование

- sklearn.model_selection.train_test_split
- sklearn.pipeline.Pipeline
- xgboost.XGBClassifier
- sklearn.linear_model.LogisticRegression
- catboost.CatBoostClassifier
- sklearn.svm.SVC
- transformers
- torch
- tensorflow.keras.models.Model
- tensorflow.keras.layers (LSTM, Activation, Dense, Dropout, Input, Embedding, SpatialDropout1D, Flatten)
- tensorflow.keras.optimizers.Adam
- tensorflow.keras.preprocessing.text.Tokenizer
- tensorflow.keras.preprocessing.sequence
- tensorflow.keras.utils.to_categorical
- tensorflow.keras.callbacks.EarlyStopping
- tensorflow.keras.models.Sequential
- tensorflow.keras.callbacks.EarlyStopping, ModelCheckpoint

Используемые библиотеки, алгоритмы, метрики

Функции

Задача 1 - HATE

```
def clean_text(text, lemma, noise_words):  
    '''Функция на вход получает текст, на выходе выдаёт очищенный текст'''  
    text = str(text).lower() # первый шаг - все тексты приводим к нижнему регистру  
  
    text = re.sub("@[\w'._-:~!]+", '', text) # второй шаг - убираем ники пользователей твитера, т.к. обычно не несут никакой окраски  
    text = re.sub('https?://\S+|www\.\S+', '', text) # третий шаг - убираем ссылки в твитах, т.к. названия ссылок обычно не влияют на тональность  
    text = re.sub('\w*\d\w*', '', text) # четвёртый шаг - убираем "слова", внутри которых есть цифры  
  
    text = re.sub('[^\w\s^.]', '', text) # пятый шаг - убираем знаки пунктуации  
    text = re.sub('[_.]+', ' ', text)  
  
    text = " ".join([word for word in text.split(' ') if word not in noise_words]) # шестой шаг - отбираем только НЕстоп-слова  
  
    text = " ".join([word.lemma_ for word in lemma(text)]) # седьмой шаг - лемматизация при помощи spacy  
    text = re.sub('[\s]+', ' ', text) # восьмой шаг - заменяем любой пробельный символ (табуляция, конец строки и т.п.) на пробел  
  
    return text
```

```
def clean_text_for_BERT(text):  
    '''Функция на вход получает текст, на выходе выдаёт очищенный текст для BERT'''  
  
    text = re.sub("@[\w'._-:~!]+", 'USER_NAME_TAG', text) # второй шаг - убираем ники пользователей твитера, т.к. обычно не несут никакой окраски  
    text = re.sub('https?://\S+|www\.\S+', 'URL_TAG', text) # третий шаг - убираем ссылки в твитах, т.к. названия ссылок обычно не влияют на тональность  
  
    return text
```

Задача 2 - TROLL

```
m = pymorphy2.MorphAnalyzer()  
def lemmatize(text, mystem=m):  
    try:  
        return ' '.join((m.parse(t)[0].normal_form for t in text.split(' ')))  
    except:  
        return " "
```

```
def clean_text_2(text, lemmatize, noise_wrods_2):  
    '''Функция на вход получает текст, на выходе выдаёт очищенный текст'''  
    text = str(text).lower() # первый шаг - все тексты приводим к нижнему регистру  
  
    text = re.sub("@[\w'._-:~!]+", '', text) # второй шаг - убираем ники пользователей твитера, т.к. обычно не несут никакой окраски  
    text = re.sub('https?://\S+|www\.\S+', '', text) # третий шаг - убираем ссылки в твитах, т.к. названия ссылок обычно не влияют на тональность  
    text = re.sub('\w*\d\w*', '', text) # четвёртый шаг - убираем "слова", внутри которых есть цифры  
  
    text = re.sub('[^\w\s^.]', '', text) # пятый шаг - убираем знаки пунктуации  
    text = re.sub('[_.]+', ' ', text)  
  
    #стоп-слова не всех языков  
    text = " ".join([word for word in text.split(' ') if word not in noise_wrods_2]) # шестой шаг - отбираем только НЕстоп-слова  
  
    text = lemmatize(text) # седьмой шаг - лемматизация при помощи spacy  
    text = re.sub('[\s]+', ' ', text) # восьмой шаг - заменяем любой пробельный символ (табуляция, конец строки и т.п.) на пробел  
  
    return text
```

Используемые библиотеки, алгоритмы, метрики

Алгоритмы

Задача 1 - HATE

Векторизация

- CountVectorizer
- TfidfVectorizer

Задача 2 - TROLL

Моделирование

- LogisticRegression
- SVC
- XGBClassifier
- CatBoostClassifier
- tensorflow.keras.models.Model(LSTM)
- transformers (Bert) + LogisticRegression
- transformers (Bert) + SVC

Используемые библиотеки, алгоритмы, метрики

Метрики

Задача 1 - HATE

Метрики

- `sklearn.metrics.classification_report`

Задача 2 - TROLL

1

Recall

2

Precision

3

F1-score

4

Accuracy

Итоги решений

5

Итоги решений

Задача 1 - HATE

CountVectorizer + LogisticRegression

```
vec = CountVectorizer(ngram_range=(1, 1))
vec.fit(df_2['tweet'])
bow = vec.transform(X_train)

clf = LogisticRegression(random_state=42, solver='liblinear', class_weight = 'balanced')
clf.fit(bow, y_train)
pred = clf.predict(vec.transform(X_test))
print(classification_report(pred, y_test))
```

	precision	recall	f1-score	support
0	0.60	0.68	0.64	247
1	0.87	0.88	0.88	816
2	0.96	0.90	0.92	856
accuracy			0.86	1919
macro avg	0.81	0.82	0.81	1919
weighted avg	0.87	0.86	0.87	1919

CountVectorizer + LogisticRegression

```
pipe3 = Pipeline([
    ('CountVectChar', CountVectorizer(analyzer='char', ngram_range=(1, 7))),
    ('LogReg', LogisticRegression(random_state=42, solver='liblinear',
                                   class_weight = 'balanced'))
])
pipe3.fit(X_train, y_train)
y_pred3 = pipe3.predict(X_test)
print(classification_report(y_pred3, y_test))
```

	precision	recall	f1-score	support
0	0.62	0.64	0.63	265
1	0.82	0.87	0.85	743
2	0.96	0.90	0.93	911
accuracy			0.85	1919
macro avg	0.80	0.80	0.80	1919
weighted avg	0.86	0.85	0.85	1919

TfidfVectorizer + LogisticRegression

```
pipe = Pipeline([
    ('tf-idf', TfidfVectorizer()),
    ('LogReg', LogisticRegression(random_state=42,
                                   solver='liblinear',
                                   class_weight = 'balanced'))
])
pipe.fit(X_train, y_train)
y_pred = pipe.predict(X_test)
print(classification_report(y_pred, y_test))
```

	precision	recall	f1-score	support
0	0.63	0.68	0.65	253
1	0.83	0.90	0.87	724
2	0.97	0.88	0.92	942
accuracy			0.86	1919
macro avg	0.81	0.82	0.81	1919
weighted avg	0.87	0.86	0.86	1919

CountVectorizer + LogisticRegression

```
pipe4 = Pipeline([
    ('CountVectChar', CountVectorizer(analyzer='char', ngram_range=(1, 9))),
    ('LogReg', LogisticRegression(random_state=42, solver='liblinear',
                                   class_weight = 'balanced'))
])
pipe4.fit(X_train, y_train)
y_pred4 = pipe4.predict(X_test)
print(classification_report(y_pred4, y_test))
```

	precision	recall	f1-score	support
0	0.61	0.66	0.63	256
1	0.83	0.87	0.85	750
2	0.96	0.90	0.93	913
accuracy			0.86	1919
macro avg	0.80	0.81	0.80	1919
weighted avg	0.86	0.86	0.86	1919

Итоги решений

Задача 1 - HATE

CountVectorizer + XGBClassifier

```
pipe5 = Pipeline([
    ('CountVectChar', CountVectorizer(tokenizer=word_tokenize,
                                      ngram_range=(1, 1))),
    ('XGB', XGBClassifier(objective = 'multi:softprob' ,
                          use_label_encoder=False,
                          eval_metric='mlogloss'))
])
pipe5.fit(X_train, y_train)
y_pred5 = pipe5.predict(X_test)
print(classification_report(y_pred5, y_test))
```

	precision	recall	f1-score	support
0	0.57	0.74	0.64	210
1	0.87	0.88	0.87	785
2	0.97	0.90	0.93	924
accuracy			0.87	1919
macro avg	0.80	0.84	0.82	1919
weighted avg	0.89	0.87	0.88	1919

CountVectorizer + CatBoostClassifier

```
pipe7 = Pipeline([
    ('CountVectChar', CountVectorizer(ngram_range=(1, 1))),
    ('CBC', CatBoostClassifier(learning_rate=0.6, depth=4,
                               loss_function='MultiClass'))
])
pipe7.fit(X_train, y_train)
y_pred7 = pipe7.predict(X_test)
print(classification_report(y_pred7, y_test))
```

	precision	recall	f1-score	support
0	0.55	0.75	0.64	210
1	0.89	0.89	0.89	809
2	0.98	0.89	0.93	900
accuracy			0.88	1919
macro avg	0.81	0.84	0.82	1919
weighted avg	0.89	0.88	0.88	1919

TfidfVectorizer + XGBClassifier

```
pipe6 = Pipeline([
    ('tf-idf', TfidfVectorizer()),
    ('XGB', XGBClassifier(booster='gbtree', objective = 'multi:softprob' ,
                          use_label_encoder=False, eval_metric='mlogloss'))
])
pipe6.fit(X_train, y_train)
y_pred6 = pipe6.predict(X_test)
print(classification_report(y_pred6, y_test))
```

	precision	recall	f1-score	support
0	0.57	0.71	0.63	222
1	0.87	0.88	0.87	781
2	0.96	0.90	0.93	916
accuracy			0.87	1919
macro avg	0.80	0.83	0.81	1919
weighted avg	0.88	0.87	0.87	1919

TfidfVectorizer + CatBoostClassifier

```
pipe8 = Pipeline([
    ('tf-idf', TfidfVectorizer(tokenizer=word_tokenize)),
    ('CBC', CatBoostClassifier(learning_rate=0.6, depth=4,
                               loss_function='MultiClass'))
])
pipe8.fit(X_train, y_train)
y_pred8 = pipe8.predict(X_test)
print(classification_report(y_pred8, y_test))
```

	precision	recall	f1-score	support
0	0.52	0.74	0.61	197
1	0.88	0.87	0.88	824
2	0.97	0.89	0.93	898
accuracy			0.87	1919
macro avg	0.79	0.83	0.80	1919
weighted avg	0.89	0.87	0.87	1919

Итоги решений

Задача 1 - HATE

TfidfVectorizer + LGBMClassifier

```
pipe0 = Pipeline([
    ('tf-idf', TfidfVectorizer(tokenizer=word_tokenize)),
    ('LGBMClass', LGBMClassifier())
])
pipe0.fit(X_train, y_train)
y_pred0 = pipe0.predict(X_test)
print(classification_report(y_pred0, y_test))
```

	precision	recall	f1-score	support
0	0.49	0.72	0.58	191
1	0.90	0.87	0.88	868
2	0.95	0.89	0.92	860
accuracy			0.86	1919
macro avg	0.78	0.83	0.80	1919
weighted avg	0.88	0.86	0.87	1919

TfidfVectorizer + LightGBM + SMOTE

```
vec_10_1 = TfidfVectorizer()
vec_10_1.fit(df_2['tweet'])
bow_10_1 = vec_10_1.transform(X_train)
```

```
sm_1 = SMOTE (#sampling_strategy = 0.9,
             random_state=0,
             k_neighbors=25)
X_train_res_1, y_train_res_1 = sm_1.fit_resample(bow_10_1, y_train)
```

```
pipe0_1 = LGBMClassifier()
pipe0_1.fit(X_train_res_1, y_train_res_1)
y_pred0_1 = pipe0_1.predict(vec_10_1.transform(X_test))
print(classification_report(y_pred0_1, y_test))
```

	precision	recall	f1-score	support
0	0.60	0.66	0.63	276
1	0.84	0.87	0.85	758
2	0.96	0.91	0.94	885
accuracy			0.86	1919
macro avg	0.80	0.81	0.81	1919
weighted avg	0.86	0.86	0.86	1919

tf.keras + Tokenizer + LSTM

```
max_words = 50000
max_len = 300
tokenizer = Tokenizer(num_words=max_words)
tokenizer.fit_on_texts(X_train)
sequences = tokenizer.texts_to_sequences(X_train)
sequences_matrix = sequence.pad_sequences(sequences, maxlen=max_len)
```

```
model_9 = Sequential()
model_9.add(Embedding(max_words, 100, input_length=max_len))
model_9.add(SpatialDropout1D(0.2))
model_9.add(LSTM(100, dropout=0.2, recurrent_dropout=0.2))
model_9.add(Dense(1, activation='relu'))
model_9.add(Dense(1, activation='sigmoid'))
model_9.summary()
model_9.compile(loss='categorical_crossentropy', optimizer=Adam(), metrics=['accuracy'])
```

```
stop = EarlyStopping(
    monitor='val_accuracy',
    mode='max',
    patience=3
)
```

```
checkpoint = ModelCheckpoint(
    filepath='.',
    save_weights_only=True,
    monitor='val_accuracy',
    mode='max',
    save_best_only=True
)
```

```
history = model_9.fit(sequences_matrix, y_train, batch_size=1024, epochs=10,
                    validation_split=0.2, callbacks=[stop, checkpoint])
```

```
Epoch 1/10
16/16 [=====] - 25s 1s/step - loss: 0.0000e+00 - accuracy: 0.7346 - val_loss: 0.0000e+00 - val_accuracy: 0.7736
Epoch 2/10
16/16 [=====] - 22s 1s/step - loss: 0.0000e+00 - accuracy: 0.7754 - val_loss: 0.0000e+00 - val_accuracy: 0.7736
Epoch 3/10
16/16 [=====] - 22s 1s/step - loss: 0.0000e+00 - accuracy: 0.7754 - val_loss: 0.0000e+00 - val_accuracy: 0.7736
Epoch 4/10
16/16 [=====] - 22s 1s/step - loss: 0.0000e+00 - accuracy: 0.7754 - val_loss: 0.0000e+00 - val_accuracy: 0.7736
```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 300, 100)	5000000
spatial_dropout1d (SpatialDropout1D)	(None, 300, 100)	0
lstm (LSTM)	(None, 100)	80400
dense (Dense)	(None, 1)	101
dense_1 (Dense)	(None, 1)	2
Total params: 5,080,503		
Trainable params: 5,080,503		
Non-trainable params: 0		

Итоги решений

Задача 1 - HATE

SMOTE + CountVectorizer + SVC

```
sm = SMOTE (#sampling_strategy = 0.9,
           random_state=0,
           k_neighbors=4)
X_train_res, y_train_res = sm.fit_resample(bow_10, y_train_10)

print('\tДО балансировки \tПОСЛЕ балансировки ')
print('класс 2 : \t{}\t{}\t{}\t{}'.format(sum(y_train_10==2), sum(y_train_res==2)))
print('класс 1 : \t{}\t{}\t{}\t{}'.format(sum(y_train_10==1), sum(y_train_res==1)))
print('класс 0 : \t{}\t{}\t{}\t{}'.format(sum(y_train_10==0), sum(y_train_res==0)))
print('y : \t{}\t{}\t{}\t{}'.format(y_train_10.shape, y_train_res.shape))
print('X : \t{}\t{}\t{}\t{}'.format(bow_10.shape, X_train_res.shape))

model_10 = SVC()
model_10.fit(X_train_res, y_train_res)
pred_10 = model_10.predict(vec_10.transform(X_test_10))
print(classification_report(pred_10, y_test_10))
```

	ДО балансировки	ПОСЛЕ балансировки
класс 2 :	3368	3368
класс 1 :	3159	3368
класс 0 :	1147	3368
y :	(7674,)	(10104,)
X :	(7674, 11677)	(10104, 11677)

	precision	recall	f1-score	support
0	0.51	0.62	0.56	240
1	0.82	0.89	0.85	716
2	0.95	0.84	0.89	963
accuracy			0.83	1919
macro avg	0.76	0.78	0.77	1919
weighted avg	0.85	0.83	0.84	1919

SMOTE + CountVectorizer + LogReg - 15000 и 5000

```
clf_00 = LogisticRegression(random_state=42, solver='liblinear')
clf_00.fit(X_train_res, y_train_res)
pred_00 = clf_00.predict(vec_10.transform(X_test_10))
print(classification_report(pred_00, y_test_10))
```

	precision	recall	f1-score	support
0	0.47	0.32	0.38	434
1	0.89	0.94	0.92	3644
2	0.87	0.82	0.84	879
accuracy			0.86	4957
macro avg	0.74	0.69	0.71	4957
weighted avg	0.85	0.86	0.86	4957

```
clf_00 = LogisticRegression(random_state=42, solver='liblinear')
clf_00.fit(X_train_res, y_train_res)
pred_00 = clf_00.predict(vec_10.transform(X_test_10))
print(classification_report(pred_00, y_test_10))
```

	precision	recall	f1-score	support
0	0.62	0.57	0.59	318
1	0.83	0.89	0.86	731
2	0.91	0.90	0.91	870
accuracy			0.84	1919
macro avg	0.79	0.78	0.79	1919
weighted avg	0.84	0.84	0.84	1919

Итоги решений

Задача 1 - HATE

BERT + LogReg - 3000

```
lr_clf_11 = LogisticRegression(class_weight = 'balanced')
lr_clf_11.fit(train_features, train_labels)
print(classification_report(lr_clf_11.predict(test_features), test_labels))
```

	precision	recall	f1-score	support
0	0.53	0.24	0.33	171
1	0.76	0.94	0.84	604
2	0.78	0.60	0.68	225
accuracy			0.74	1000
macro avg	0.69	0.59	0.61	1000
weighted avg	0.72	0.74	0.71	1000

BERT + LogReg - 1000 - 0.2 + SMOTE

```
lr_clf_11 = LogisticRegression(class_weight = 'balanced')
lr_clf_11.fit(X_train_res, y_train_res)
print(classification_report(lr_clf_11.predict(test_features), test_labels))
```

	precision	recall	f1-score	support
0	0.39	0.39	0.39	18
1	0.86	0.88	0.87	137
2	0.72	0.69	0.70	45
accuracy			0.79	200
macro avg	0.66	0.65	0.65	200
weighted avg	0.79	0.79	0.79	200

BERT + LogReg - 1000

```
lr_clf_11 = LogisticRegression(class_weight = 'balanced')
lr_clf_11.fit(train_features, train_labels)
print(classification_report(lr_clf_11.predict(test_features), test_labels))
```

	precision	recall	f1-score	support
0	0.45	0.36	0.40	14
1	0.86	0.93	0.89	145
2	0.72	0.56	0.63	41
accuracy			0.81	200
macro avg	0.68	0.62	0.64	200
weighted avg	0.80	0.81	0.81	200

BERT + LogReg - 1000 - 0.35 + SMOTE

```
lr_clf_11 = LogisticRegression(class_weight = 'balanced')
lr_clf_11.fit(X_train_res, y_train_res)
print(classification_report(lr_clf_11.predict(test_features), test_labels))
```

	precision	recall	f1-score	support
0	0.33	0.40	0.36	20
1	0.89	0.90	0.89	259
2	0.69	0.61	0.65	71
accuracy			0.81	350
macro avg	0.64	0.64	0.64	350
weighted avg	0.82	0.81	0.81	350

Итоги решений

Задача 1 - HATE

BERT + SVC- 3000

```
lr_clf_12 = SVC(class_weight = 'balanced')
lr_clf_12.fit(train_features, train_labels)
print(classification_report(lr_clf_12.predict(test_features), test_labels))
```

	precision	recall	f1-score	support
0	0.55	0.23	0.32	189
1	0.72	0.94	0.82	575
2	0.79	0.58	0.67	236
accuracy			0.72	1000
macro avg	0.69	0.58	0.60	1000
weighted avg	0.71	0.72	0.69	1000

BERT + SVC- 1000 - 0.2/0.35

```
lr_clf_12 = SVC(class_weight = 'balanced')
lr_clf_12.fit(train_features, train_labels)
print(classification_report(lr_clf_12.predict(test_features), test_labels))
```

	precision	recall	f1-score	support
0	0.36	0.13	0.20	30
1	0.76	0.94	0.84	126
2	0.66	0.48	0.55	44
accuracy			0.72	200
macro avg	0.59	0.52	0.53	200
weighted avg	0.68	0.72	0.68	200

```
lr_clf_12 = SVC(class_weight = 'balanced')
lr_clf_12.fit(train_features, train_labels)
print(classification_report(lr_clf_12.predict(test_features), test_labels))
```

	precision	recall	f1-score	support
0	0.57	0.15	0.24	52
1	0.77	0.96	0.85	221
2	0.72	0.56	0.63	77
accuracy			0.75	350
macro avg	0.69	0.56	0.57	350
weighted avg	0.73	0.75	0.71	350

BERT + SVC- 1000 - 0.2/0.35 + SMOTE

```
lr_clf_12 = SVC(class_weight = 'balanced')
lr_clf_12.fit(X_train_res, y_train_res)
print(classification_report(lr_clf_12.predict(test_features), test_labels))
```

	precision	recall	f1-score	support
0	0.61	0.38	0.47	29
1	0.81	0.94	0.87	119
2	0.77	0.63	0.69	52
accuracy			0.78	200
macro avg	0.73	0.65	0.68	200
weighted avg	0.77	0.78	0.77	200


```
lr_clf_12 = SVC(class_weight = 'balanced')
lr_clf_12.fit(X_train_res, y_train_res)
print(classification_report(lr_clf_12.predict(test_features), test_labels))
```

	precision	recall	f1-score	support
0	0.54	0.30	0.39	43
1	0.80	0.94	0.86	223
2	0.73	0.54	0.62	84
accuracy			0.77	350
macro avg	0.69	0.59	0.62	350
weighted avg	0.75	0.77	0.75	350

Итоги решений

Задача 1 - NATE

Визуализация лучших алгоритмов Задача 1



	TfLr	CCat	TfLgSm	CLrSm5000	BERTLr1000	BERTSvc1000
0	0.65	0.64	0.63	0.59	0.40	0.47
1	0.87	0.89	0.85	0.86	0.89	0.87
2	0.92	0.93	0.94	0.91	0.63	0.69
accuracy	0.86	0.88	0.86	0.84	0.81	0.78
	✗	✓	✗	✗	✗	✗
	②	①	③	④	⑤	⑥

Итоги решений

Задача 2 - TROLL

CountVectorizer + LogisticRegression

```
vec_2 = CountVectorizer(ngram_range=(1, 1))
vec_2.fit(data_total_4_English['content'].values.astype('U'))
bow_2 = vec_2.transform(X_train_E_2)

clf_2_1 = LogisticRegression(random_state=42, solver='liblinear',
                             class_weight = 'balanced')
clf_2_1.fit(bow_2, y_train_E_2)
pred_2_1 = clf_2_1.predict(vec_2.transform(X_test_E_2))
print(classification_report(pred_2_1, y_test_E_2))
```

	precision	recall	f1-score	support
0	0.95	0.94	0.94	10628
1	0.68	0.28	0.40	2544
2	0.83	0.80	0.81	22712
3	0.67	0.69	0.68	38740
4	0.88	0.86	0.87	57489
5	0.59	0.29	0.39	5420
6	0.72	0.83	0.77	58206
7	0.25	0.04	0.07	4261
accuracy			0.78	200000
macro avg	0.69	0.59	0.62	200000
weighted avg	0.77	0.78	0.77	200000

TfidfVectorizer + LogisticRegression

```
pipe_2_2 = Pipeline([
    ('tf-idf', TfidfVectorizer()),
    ('LogReg', LogisticRegression(random_state=42,
                                   solver='liblinear',
                                   class_weight = 'balanced'))
])

pipe_2_2.fit(X_train_E_2, y_train_E_2)
y_pred_2_2 = pipe_2_2.predict(X_test_E_2)
print(classification_report(y_pred_2_2, y_test_E_2))
```

	precision	recall	f1-score	support
0	0.96	0.92	0.94	10927
1	0.68	0.29	0.41	2464
2	0.81	0.85	0.83	20844
3	0.67	0.67	0.67	39872
4	0.88	0.85	0.87	58209
5	0.60	0.27	0.38	5995
6	0.71	0.83	0.76	57654
7	0.25	0.04	0.07	4035
accuracy			0.77	200000
macro avg	0.70	0.59	0.62	200000
weighted avg	0.76	0.77	0.76	200000

Итоги решений

Задача 2 - TROLL

CountVectorizer + XGBClassifier

```
pipe5_2 = Pipeline([
    ('CountVectChar', CountVectorizer(ngram_range=(1, 1))),
    ('XGB', XGBClassifier(objective = 'multi:softprob' ,
                          use_label_encoder=False,
                          eval_metric='mlogloss'))
])
```

```
pipe5_2.fit(X_train_E_2, y_train_E_2)
y_pred5_2 = pipe5_2.predict(X_test_E_2)
print(classification_report(y_pred5_2, y_test_E_2))
```

	precision	recall	f1-score	support
0	0.77	0.94	0.85	8531
1	0.50	0.82	0.62	643
2	0.57	0.98	0.72	12713
3	0.42	0.71	0.53	23654
4	0.76	0.84	0.80	51063
5	0.17	0.85	0.28	537
6	0.86	0.56	0.68	102833
7	0.03	0.65	0.05	26
accuracy			0.70	200000
macro avg	0.51	0.79	0.57	200000
weighted avg	0.76	0.70	0.70	200000

CountVectorizer + XGBClassifier

```
pipe7_2 = Pipeline([
    ('CountVectChar', CountVectorizer(ngram_range=(1, 1))),
    ('CBC', CatBoostClassifier( learning_rate=1, depth=2,
                               loss_function='MultiClass'))
])
```

```
pipe7_2.fit(X_train_E_2, y_train_E_2)
y_pred7_2 = pipe7_2.predict(X_test_E_2)
print(classification_report(y_pred7_2, y_test_E_2))
```

	precision	recall	f1-score	support
0	0.66	0.95	0.78	7365
1	0.32	0.89	0.47	383
2	0.37	1.00	0.54	8195
3	0.25	0.60	0.36	16715
4	0.64	0.82	0.72	43865
5	0.00	0.25	0.00	4
6	0.87	0.47	0.61	123473
7	0.00	0.00	0.00	0
accuracy			0.60	200000
macro avg	0.39	0.62	0.44	200000
weighted avg	0.74	0.60	0.62	200000

TfidfVectorizer + XGBClassifier

```
pipe6_2 = Pipeline([
    ('tf-idf', TfidfVectorizer()),
    ('XGB', XGBClassifier(objective = 'multi:softprob' ,
                          use_label_encoder=False,
                          eval_metric='mlogloss'))
])
```

```
pipe6_2.fit(X_train_E_2, y_train_E_2)
y_pred6_2 = pipe6_2.predict(X_test_E_2)
print(classification_report(y_pred6_2, y_test_E_2))
```

	precision	recall	f1-score	support
0	0.77	0.95	0.85	8498
1	0.49	0.81	0.61	646
2	0.57	0.99	0.72	12718
3	0.42	0.72	0.53	23034
4	0.77	0.84	0.80	51026
5	0.18	0.80	0.29	600
6	0.87	0.56	0.68	103444
7	0.03	0.56	0.05	34
accuracy			0.70	200000
macro avg	0.51	0.78	0.57	200000
weighted avg	0.76	0.70	0.70	200000

Итоги решений

Задача 2 - TROLL

TfidfVectorizer+ LGBMClassifier

```
pipe0 = Pipeline([
    ('tf-idf', TfidfVectorizer()),
    ('LGBMClass', LGBMClassifier())
])
pipe0.fit(X_train_E_2, y_train_E_2)
y_pred0 = pipe0.predict(X_test_E_2)
print(classification_report(y_pred0, y_test_E_2))
```

	precision	recall	f1-score	support
0	0.82	0.94	0.87	9129
1	0.40	0.63	0.49	670
2	0.63	0.98	0.76	14041
3	0.51	0.67	0.58	30104
4	0.82	0.83	0.83	55709
5	0.22	0.75	0.34	798
6	0.83	0.62	0.71	89419
7	0.06	0.32	0.10	130
accuracy			0.73	200000
macro avg	0.54	0.72	0.59	200000
weighted avg	0.76	0.73	0.73	200000

TfidfVectorizer + LGBMClassifier + SMOTE

```
vec_10_1 = TfidfVectorizer()
vec_10_1.fit(data_total_4_English['content'].values.astype('U'))
bow_10_1 = vec_10_1.transform(X_train_E_2)

sm_1 = SMOTE (#sampling_strategy = 0.9,
              random_state=0,
              k_neighbors=100)
X_train_res_1, y_train_res_1 = sm_1.fit_resample(bow_10_1, y_train_E_2)

pipe0_1 = LGBMClassifier()

pipe0_1.fit(X_train_res_1, y_train_res_1)
y_pred0_1 = pipe0_1.predict(vec_10_1.transform(X_test_E_2))
print(classification_report(y_pred0_1, y_test_E_2))
```

	precision	recall	f1-score	support
0	0.86	0.87	0.86	10437
1	0.64	0.38	0.48	1797
2	0.66	0.96	0.78	15032
3	0.62	0.52	0.56	47534
4	0.81	0.84	0.82	54046
5	0.51	0.16	0.25	8533
6	0.62	0.80	0.70	51678
7	0.34	0.02	0.04	10943
accuracy			0.69	200000
macro avg	0.63	0.57	0.56	200000
weighted avg	0.67	0.69	0.66	200000

Итоги решений

Задача 2 - TROLL

SMOTE + CountVectorizer + LogisticRegression (300 000)

```
X_train_E_2, X_test_E_2, y_train_E_2, y_test_E_2 = train_test_split(data_total_4_English['content'].values.astype('U'),
                                                                    data_total_4_English['account_category'] ,
                                                                    test_size = 0.2)
```

```
vec_10_E = CountVectorizer( ngram_range=(1, 1))
vec_10_E.fit(data_total_4_English['content'].values.astype('U'))
bow_10_E = vec_10_E.transform(X_train_E_2)
```

```
sm = SMOTE (#sampling_strategy = 0.9,
           random_state=0,
           k_neighbors=25)
X_train_res_E, y_train_res_E = sm.fit_resample(bow_10_E, y_train_E_2)
```

```
print('\t\tДО балансировки \tПОСЛЕ балансировки ')
print('y : \t\t{}\t\t{}'.format(y_train_E_2.shape, y_train_res_E.shape))
print('X : \t\t{}\t\t{}'.format(bow_10_E.shape, X_train_res_E.shape))
```


```
          ДО балансировки      ПОСЛЕ балансировки
y :      (800000,)             (2137584,)
X :      (800000, 212850)       (2137584, 212850)
```

```
: model_8_1 = LogisticRegression(random_state=42, solver='liblinear')
model_8_1.fit(X_train_res_E, y_train_res_E)
pred_8_1 = model_8_1.predict(vec_10_E.transform(X_test_E_2))
print(classification_report(pred_8_1, y_test_E_2))
```

	precision	recall	f1-score	support
0	0.93	0.94	0.94	10575
1	0.58	0.33	0.42	1768
2	0.83	0.81	0.82	22708
3	0.66	0.69	0.68	38335
4	0.89	0.85	0.87	58710
5	0.45	0.29	0.35	4008
6	0.74	0.82	0.77	60677
7	0.14	0.03	0.05	3219
accuracy			0.78	200000
macro avg	0.65	0.60	0.61	200000
weighted avg	0.77	0.78	0.78	200000

Итоги решений

Задача 2 - TROLL



	CLr	TfXg	TfLg	CLrSm300000
0	0.94	0.85	0.87	0.94
1	0.40	0.61	0.49	0.42
2	0.81	0.72	0.76	0.82
3	0.68	0.53	0.58	0.68
4	0.87	0.80	0.83	0.87
5	0.39	0.29	0.34	0.35
6	0.77	0.68	0.71	0.77
7	0.07	0.05	0.10	0.05
accuracy	0.78	0.70	0.73	0.78
	✓	✗	✗	✗
	①	③	④	②

Выводы



6

Выводы

1. Лучше всего для задач определения тональности текста подходит алгоритм:
 - a. LogisticRegression
2. Ансамблевые модели показывают не лучшие результаты на большом количестве данных, чем логистическая регрессия.
3. В дальнейшем, повысить качество можно:
 - a. путём использованием всех столбцов данных
 - b. настройкой гиперпараметров при помощи GridSearchCV, RandomizedSearchCV
 - c. запуском обучения, используя большее количество ресурсов

Список источников



7

СПИСОК ИСТОЧНИКОВ

- <http://neerc.ifmo.ru/wiki/>
- <http://neerc.ifmo.ru/wiki/index.php>
- <https://proglib.io/p/analiz-tonalnosti-teksta-proshloe-nastoyashchee-i-budushchee-2020-11-30>



Николай Павлов

Junior Data Scientist

Аккаунты в соцсетях



vk.com/id51647681



t.me/NGPavlov



github.com/PNikolayG

Почта



pavlovnikg@gmail.com

Резюме



https://docs.google.com/document/d/1OUWezQMI4xVgL7o3v7PJ_VxnIhk5cvKbFqJiRjLK4kA/edit?usp=sharing

*“Буду Вам очень благодарен за
рекомендацию!”*

Анализ эмоциональной окраски текста

Спасибо за внимание!

Николай Павлов
группа DS-27



pavlovnikg@gmail.com



t.me/NGPavlov